

Generative AI for Finance and Fintech

Learner Guide

Concepts, detailed procedures, lab evidence and responsible finance-AI practice

Course code	TGS-2026065050
TSC	ACC-ICT-5004-1.1 · Digital Technology Adoption and Innovation
Version	v5.0 · 6 September 2026
Organisation	Tertiary Infotech Academy Pte Ltd
Duration	4 days · 30 classroom hours + 2 assessment hours

Controlled learning document

Document Version Control Record

Document	Version	Date	Change summary
Learner Guide	5.0	6 September 2026	Complete redesign aligned to TGS-2026065050, current Microsoft copilots, finance AI use cases, labs, security and assessment evidence.

Table of Contents

Open in Word and update this field if page numbers change.

How to use this guide

Slides versus guide

The trainer deck is concept-led. Detailed click-by-click and evidence procedures are intentionally contained here and in the individual lab folders.

Learning outcomes

- LO1: Support Artificial Intelligence (AI) and Generative AI innovations to champion organisation-wide innovation.
- LO2: Evaluate and implement GAI technologies for the organisation.
- LO3: Review and adopt new GAI technologies by reviewing use cases in finance operations.
- LO4: Review AI ethics and governance in organisation processes for proactive risk monitoring and safety.
- LO5: Review emerging trends in GAI and strategise cost-control measures to support innovation initiatives.

Course map

Topic	Focus	Criteria
1	Generative AI, Agentic AI and AI Agents in Finance	K2 · A1
2	Excel Copilot for Financial Analysis and Dashboards	K1 · A2
3	SharePoint for Finance and Grounded Finance Agents	K1 · A3
4	Copilot Studio Workflows and Agents for Finance Automation	A3
5	Multi-Agent Orchestration of Financial Work Processes	K4 · A3
6	AI Security, PDPA and Human Oversight in Finance	K3 · A4 · A5

Assessment

- Written Assessment (SAQ): 4 open-ended questions, 60 minutes, open book, K2/K1/K3/K4.
- Case Study (CS): 5 integrated tasks, 60 minutes, open book, A1-A5.
- No practice examination is included. The assessment mirrors the live legacy instrument count, mapping and timing.

Topic 1: Generative AI, Agentic AI and AI Agents in Finance

History of Generative AI, how the three types differ, finance and fintech use cases for each, and Copilot in Word, PowerPoint and the M365 agent builder

Competency focus: K2 · A1

Core concepts

Why the three terms matter

Generative AI, Agentic AI and AI Agents are three different things, and finance controls differ for each.

- Generative AI produces content on request
- Agentic AI pursues a goal across steps
- An AI Agent is the deployed unit that does it
- Confusing them produces the wrong control design

A short history of Generative AI

Today's Copilot rests on seventy years of steady narrowing from rules to learned representation.

- 1950s-80s: symbolic rules and expert systems
- 1990s-2000s: statistical machine learning
- 2017: the transformer architecture
- 2022-2026: assistants, then Copilot agents in the workplace

From ELIZA to Copilot

Each generation moved more of the pattern from the programmer into the data.

- ELIZA scripted a conversation
- Credit scorecards learned from labelled outcomes
- Language models learned from unlabelled text
- Agents now select their own tools and knowledge

What Generative AI actually does

It predicts likely content, which explains both its fluency and its errors.

- Drafts narrative, summaries and code
- Transforms text, tables and slides
- Has no inherent access to your ledger
- Fluent output is not verified output

What makes AI Agentic

Agency is the ability to plan and act, not the ability to write well.

- Decomposes a goal into steps
- Chooses tools and knowledge at run time
- Observes results and adapts
- Persists context across a task

What an AI Agent is made of

An agent is a governed assembly of instructions, knowledge, tools, identity and memory.

- Instructions define purpose and boundary
- Knowledge grounds the answer
- Tools perform permitted actions
- Identity limits what it may reach

The three compared

Autonomy rises across the three, and so must the control.

- Generative: you prompt, it drafts, you own the output
- Agentic: it plans and acts within a boundary you set
- Agent: a named, permissioned, monitored deployment
- Evidence and approval requirements rise in step

GenAI use cases in finance

Generative AI earns its place where drafting and explanation dominate.

- Management commentary and variance narrative
- Board pack and investor summary drafting
- Credit memo and policy summarisation
- Meeting and audit evidence summarisation

Agentic AI use cases in finance

Agentic AI suits repeatable multi-step processes with clear rules.

- Month-end close task orchestration
- Invoice and payment exception routing
- Reconciliation preparation and follow-up
- Collections triage with drafted outreach

AI Agent use cases in fintech

Fintech deploys agents where volume and latency defeat manual work.

- Customer onboarding and KYC pre-checks
- Payment fraud triage and enrichment
- Robo-advisory and portfolio explanation
- RegTech obligation tracking and filing support

Copilot in Word for finance

Word Copilot drafts the report; the finance professional owns every number in it.

- Draft from an approved data source or prompt
- Reference an existing workbook or document
- Rewrite for tone, length and audience
- Verify every figure against the source

Copilot in PowerPoint for finance

PowerPoint Copilot turns an approved document into a first-draft deck.

- Create a deck from an existing Word file
- Generate speaker notes and summaries

- Restructure and add slides on request
- Correct the story; the model does not know materiality

The Word-to-PowerPoint pipeline

Drafting the report first and generating the deck from it keeps one version of the truth.

- Approve the workbook figures first
- Draft the report in Word from that source
- Generate the deck from the approved report
- Reconcile every repeated number

Prompt contracts for finance

A strong finance prompt names role, task, source boundary, controls and output shape.

- State the finance objective and period
- Name the approved source and as-of date
- Require caveats and reconciliation
- Specify tables, units and review status

Human accountability

AI prepares and recommends; an accountable person approves.

- Define preparer and reviewer roles
- Set monetary and risk thresholds
- Require evidence for any override
- Escalate ambiguity rather than invent

Finance use-case catalogue

Use case	Operating pattern	Value	Risk / control
Management report drafting	Draft the monthly report in Word from approved figures	Shorter reporting cycle	Unverified numbers reaching the board / Every figure traced to the approved workbook before issue
Board deck creation	Generate a PowerPoint deck from the approved Word report	Consistent single story	Deck and report drifting apart / Generate from the approved document, never from memory
Credit memo drafting	Assemble borrower facts and analyst rationale	Analyst productivity	Bias and unsupported conclusions / Reason codes, citations and credit approval
Regulatory summarisation	Extract obligations and draft a review checklist	Faster interpretation	Hallucinated requirement / Primary-source citations and legal review
Finance policy	Answer staff policy	Fewer support queries	Outdated or invented

helpdesk	questions with citations		advice / Grounded on approved documents with effective dates
Month-end task orchestration	Sequence close tasks, chase evidence, surface exceptions	Predictable close	Autonomous sign-off / Read-only by default, explicit human approval
Invoice exception routing	Detect missing fields and route to the right owner	Lower handling time	Duplicate or fraudulent payment / Three-way match and reviewer approval
Customer onboarding checks	Run KYC pre-checks and assemble a reviewer pack	Faster onboarding	Wrongful rejection and bias / Human adjudication and appeal route
Payment fraud triage	Enrich and prioritise suspicious transactions	Lower fraud loss	False positives and customer friction / Thresholds tuned to investigator capacity
Robo-advisory explanation	Explain a portfolio recommendation in plain language	Better client understanding	Unsuitable advice / Scope limits and adviser accountability

Real-world cases

From ELIZA to Copilot

Seventy years moved the pattern from the programmer into the data: scripted rules, then statistical learning, then transformers, and now agents that select their own tools. Each step widened capability and shifted where the controls must sit.

Source: [Technology history](#)

Morgan Stanley knowledge assistant

A grounded assistant helps advisers navigate approved internal research while the adviser remains accountable for the client advice. A clear example of generative assistance without transferring accountability.

Source: [Wealth management](#)

Northstar Finance Helper

The M365 agent built in Lab 3 answers a claim-approval question by citing POL-FIN-001, its effective date, the SGD 2,000.01 to 10,000 threshold and the 30-day deadline. Built with no code in about five minutes.

Source: [Course tenant](#)

Control reflection

- Which source is authoritative and current?
- Which number, formula or action needs independent review?
- What evidence proves the control ran?
- When must the user refuse, escalate or fall back?

Topic 2: Excel Copilot for Financial Analysis and Dashboards

Data readiness, Copilot-generated formulas, variance analysis, visualisation and management dashboards

Competency focus: K1 · A2

Core concepts

Copilot in Excel today

Copilot edits native Excel objects, so the result stays auditable and editable.

- Works on tables, formulas and ranges
- Creates charts and PivotTables
- Offers chat, plan and edit modes
- Leaves inspectable native artefacts behind

Workbook readiness

Reliable Copilot analysis starts with a flat, typed, labelled table.

- One row per business event
- Stable headers, no merged cells
- Real dates and numeric currency
- Unique keys and clear dimensions

Why messy data breaks Copilot

Copilot infers structure; ambiguous structure produces confident wrong answers.

- Text-formatted numbers do not sum
- Merged cells hide the real grain
- Mixed date formats disorder periods
- Blank keys silently drop rows

Convert to an Excel Table

Ctrl+T is the single highest-value action before using Copilot.

- Gives the range a name and a grain
- Enables structured references
- Makes columns explicit to Copilot
- Required by downstream connectors

Chat, plan or edit

Choose the least invasive Copilot mode that meets the task.

- Chat for read-only exploration
- Plan to review the intended approach
- Edit for bounded workbook change
- Review every change before accepting

Formula generation

A generated formula is a proposal that needs the same audit as a written one.

- Prefer transparent native functions
- Check ranges and absolute references
- Test blanks and edge cases
- Reconcile to a control total

The finance formula toolkit

A small vocabulary covers most operational analysis.

- SUMIFS for controlled aggregation
- XLOOKUP for mapped dimensions
- IFERROR for explicit exceptions
- FORECAST.LINEAR for simple projection

Budget versus actual

Variance is only meaningful when periods, scope and sign convention match.

- Compare like periods and entities
- Fix the sign convention first
- Separate volume from rate effects
- Define materiality before explaining

Variance analysis with Copilot

Copilot finds the largest movements; the analyst supplies the cause.

- Rank variances by value and percentage
- Apply a materiality threshold
- Identify the top contributors
- Reject unsupported causal claims

Visualisation choices

The chart type is a claim about the data's shape.

- Column for period comparison
- Line for trend over time
- Waterfall for bridges and drivers
- Avoid pie charts for finance detail

Dashboard design

A finance dashboard is a decision surface, not a chart collage.

- Show actual, budget and variance
- Keep time and currency units consistent
- Surface exceptions and drivers
- Link every visual to source data

PivotTables and Copilot

PivotTables give Copilot a clean aggregate to describe.

- Build from a single flat table
- Keep the period ordering explicit
- Refresh before asking for narrative
- Retain the comparison basis

Forecasting in Excel

Copilot accelerates model assembly; management still owns the assumptions.

- Start from operational drivers
- Separate base, upside and downside
- Use assumption cells, not hard-codes
- Back-test against actuals

Management commentary

A useful narrative quantifies, attributes and assigns an action.

- State magnitude and period
- Separate fact from hypothesis
- Name owner and due date
- Cite the source cell or table

Excel control checklist

Prove structure, formulas, outputs and access before sharing.

- Run a formula error scan
- Reconcile control totals
- Test sensitivity and edge cases
- Protect source and assumption ranges

Finance use-case catalogue

Use case	Operating pattern	Value	Risk / control
Data readiness triage	Find blanks, duplicates and text-formatted numbers	Analysis that can be trusted	Silent row loss / Control totals before and after cleaning
Budget versus actual analysis	Summarise variances by department and account	Shorter reporting cycle	False causality / Materiality rules and evidence-backed narrative
Formula generation and audit	Draft SUMIFS, XLOOKUP and IFERROR formulas	Faster model assembly	Wrong ranges and hidden hard-codes / Formula audit and reconciliation to control totals
Management dashboard	Build KPI tiles and linked charts	Actionable single view	Metric inconsistency / Shared definitions and one source table
Revenue bridge	Decompose growth into volume, price and mix	Clearer performance story	Incorrect attribution / Reconciled bridge and driver definitions
Rolling forecast	Extend the forecast from recent actuals and drivers	Faster reforecast	Overfitting and stale assumptions / Assumption cells,

			scenarios and back-testing
Working-capital dashboard	Visualise DSO, ageing and cash conversion	Actionable cash view	Stale as-at dates / Dated extracts and control totals
Spend anomaly review	Highlight unusual vendor, amount or timing patterns	Leakage reduction	False accusation / Risk scores used as triage only
Scenario planning	Model base, downside and upside coherently	Better decisions under uncertainty	Isolated cell edits / Scenarios driven by assumption cells
Commentary drafting	Draft variance commentary from the analysis table	Faster reporting	Unsupported causal claims / Facts and hypotheses labelled separately

Real-world cases

Finance Agent variance analysis

Microsoft describes AI identifying material period variances, contributing factors and a structured narrative from flat source data and PivotTables. The pattern only works when the source table is clean.

Source: [FP&A](#)

Finance Agent reconciliation

Automated reconciliation reports, discrepancy suggestions and audit-ready summaries in Excel, with the reviewer retaining sign-off.

Source: [Controllershship](#)

Control reflection

- Which source is authoritative and current?
- Which number, formula or action needs independent review?
- What evidence proves the control ran?
- When must the user refuse, escalate or fall back?

Topic 3: SharePoint for Finance and Grounded Finance Agents

Finance sites, lists and libraries, P&L, forecasting and sales data, and agents grounded on approved finance content

Competency focus: K1 · A3

Core concepts

Why SharePoint for finance AI

An agent can only be as governed as the source it is grounded on.

- Central approved location for finance content
- Permissions the agent inherits
- Version history and audit trail
- Native Copilot and Graph connectivity

Site design for a finance team

Structure the site around finance processes, not around file types.

- A hub for finance, spokes per process
- Separate policy from transactional data
- Restrict sensitive areas by group
- Name libraries for the process they serve

Lists versus libraries

Lists hold structured records; libraries hold documents.

- Lists for P&L lines, forecasts and sales records
- Libraries for policies, reports and packs
- Lists give typed columns and filters
- Libraries give versioning and co-authoring

Designing a finance list

Column types decide whether an agent can filter and compute.

- Use number and currency, not text
- Use date columns for periods
- Use choice columns for controlled values
- Keep one row per business fact

Finance data model for the labs

The lab tenant carries five connected finance datasets.

- General ledger and chart of accounts
- Monthly P&L by department
- Rolling forecast and assumptions
- Sales pipeline and collections status

Permissions and inheritance

The agent should never see more than the person asking.

- Inherit from the site by default
- Break inheritance only with a reason
- Use groups, never individual grants
- Review access on a schedule

Sensitivity labels

A label travels with the file and enforces handling downstream.

- Classify finance content on creation
- Restrict export and external sharing
- Labels drive DLP enforcement
- Unlabelled content is the weak point

Grounding a Copilot agent on SharePoint

Point the agent at a folder or site, not at the whole tenant.

- Add SharePoint as a knowledge source
- Use the narrowest URL that works
- Encode spaces in the folder URL
- Verify the agent returns citations

What SharePoint grounding does well

Retrieval works best on text the model can read and cite.

- Policies, procedures and memos
- Reports and board packs
- Structured lists with clear columns
- Content with clear effective dates

Where grounding disappoints

Retrieval is not calculation, and finance numbers need arithmetic.

- Do not ask an agent to total a ledger
- Complex spreadsheets retrieve poorly
- Scanned PDFs may not be readable
- Compute in Excel, explain with the agent

Citations and effective dates

A finance answer without a source and a date is not usable evidence.

- Require citations in the instructions
- Surface the effective date
- Escalate when policy conflicts
- Retire superseded documents

Agent instructions for finance

Instructions are the boundary the model works inside.

- State the finance role and scope
- Constrain to the named knowledge
- Require refusal when unsupported
- Forbid advice outside authority

Testing a grounded agent

Test the refusals as carefully as the answers.

- Positive: a question the source answers
- Negative: a question it must refuse
- Stale: a superseded policy
- Boundary: a question outside scope

Finance use-case catalogue

Use case	Operating pattern	Value	Risk / control
Finance policy hub	Centralise approved policies with owners and dates	Single source of truth	Superseded documents in use / Effective dates, versioning and retirement
Period report library	Publish P&L, ageing and reconciliation reports	Consistent reporting	Ad-hoc spreadsheets circulating / Approved library with sensitivity labels
Structured finance lists	Hold GL, P&L, forecast and sales as typed lists	Filterable, connectable data	Text-typed numbers and dates / Column types, choice fields and validation
Grounded policy agent	Answer policy questions from approved content only	Reduced support load	Oversharing and invented clauses / Authentication, source scoping, refusal design
Permission-aware retrieval	Return only what the asker may already see	Safe self-service	Cross-team data exposure / Inherited permissions and access reviews
Close evidence retrieval	Find the reconciliation evidence behind a balance	Faster audit response	Incomplete evidence / Evidence IDs and reviewer sign-off
Sensitivity labelling	Classify finance content on creation	Enforced handling downstream	Unlabelled sensitive content / Labels drive DLP and sharing limits
Agent knowledge scoping	Point an agent at a folder, not the tenant	Precise answers	Retrieval noise and oversharing / Narrowest working scope and tested refusals

Real-world cases

Finance query management

Microsoft's scenario library grounds an agent on approved finance content to answer internal process questions, preserving source permissions.

Source: [Finance operations](#)

Grounding is not calculation

During the build, an agent pointed at the Northstar site retrieved only the site pages, not the list rows. SharePoint knowledge indexes documents. The fix was to publish the finance data as reports as well as lists.

Source: [Course tenant](#)

Control reflection

- Which source is authoritative and current?
- Which number, formula or action needs independent review?
- What evidence proves the control ran?
- When must the user refuse, escalate or fall back?

Topic 4: Copilot Studio Workflows and Agents for Finance Automation

A dedicated environment, agent flows, skills and tools, SharePoint connections and human-in-the-loop approval gates

Competency focus: A3

Core concepts

Why a dedicated environment

An environment is the blast radius for everything built inside it.

- Separates course work from production
- Scopes connectors and data policies
- Enables clean reset between cohorts
- Carries its own Dataverse store

The course environment

All Topic 4 and 5 work happens in one named sandbox environment.

- Named for the course code
- Sandbox type, Asia region, SGD currency
- Dataverse enabled for agents and flows
- Open access for the training cohort

Copilot Studio building blocks

Knowledge, tools, topics and instructions form one control system.

- Knowledge answers questions
- Tools take actions
- Topics handle deterministic paths
- Instructions set the boundary

Agent flows

An agent flow is the deterministic spine an agent can call.

- Triggered by an agent or a schedule
- Steps run in a defined order
- Connectors reach SharePoint and Excel
- Published before an agent can use it

Designing a finance workflow

Automate the repeatable middle, not the judgement at either end.

- Trigger on a finance event
- Retrieve from an approved source
- Apply deterministic rules first
- Route the exception to a person

Skills and tools

Tools expand capability and risk in the same motion.

- Add a workflow as a tool
- Separate read tools from write tools
- Constrain each tool's scope
- A published workflow only is selectable

Connecting to SharePoint

The connection carries an identity, and that identity has permissions.

- Pick the site from the dropdown
- Use Get items with a filter query
- Normalise the key inside the filter
- Never widen permissions to fix a lookup

Human-in-the-loop

The pause is the control, and it must be structural.

- The Human review node always fires
- Agent-initiated help is not a control
- Define the inputs the node publishes
- Route the request through Teams

Designing the approval gate

A gate only works if the default fails safe.

- Leave the outcome blank by default
- Name the approver explicitly
- Log the draft before the gate
- Record the decision and the responder

Approval thresholds

Not every action deserves the same friction.

- Set monetary thresholds
- Escalate by risk tier, not volume
- Auto-approve only reversible actions
- Review threshold settings periodically

Testing a finance workflow

A green run does not mean the work happened.

- Read the run's node outputs
- Test the rejection path
- Test a missing or malformed input
- Confirm published, not draft

Publishing and versions

The endpoint serves the published version, not your draft.

- Publish after every change
- Check version history
- Test after publishing, not before
- Roll back through a known-good version

Finance use-case catalogue

Use case	Operating pattern	Value	Risk / control
Invoice approval workflow	Route an invoice by amount to the right approver	Fewer bottlenecks	Payment released without approval / Threshold routing and a blocking approval gate
Collections outreach approval	Draft a reminder and hold it for human approval	Scalable outreach	Message sent without review / Human review node with fail-safe default
Expense exception handling	Gather missing evidence and route exceptions	Reduced queue ageing	Tool abuse and over-permissioned actions / Whitelisted actions and amount thresholds
Reconciliation preparation flow	Pull bank and ledger rows and pre-classify matches	Faster close	Auto-clearing real exceptions / Deterministic rules first, reviewer clears
Journal request routing	Assemble a journal request pack for approval	Better evidence quality	Agent posting a journal / Structural prohibition on posting or approving
Vendor onboarding checks	Validate details and route for approval	Faster onboarding	Fraudulent vendor details / Independent verification and dual approval
Purchase requisition triage	Classify requests and apply spending thresholds	Faster approvals	Threshold bypass / Thresholds from the delegation matrix
Audit evidence assembly	Collect and package evidence for a test procedure	Reduced search time	Evidence incompleteness / Evidence IDs and auditor judgement

Real-world cases

Collections agent

An agent combines customer and ERP context with a drafted message while a person retains accountability for external contact.

Source: [Working capital](#)

The default that undid a control

A Human review node whose outcome defaults to Approve converts a gate into a rubber stamp, and the flow diagram looks identical either way. Leave the default blank so inattention fails safe.

Source: [Course tenant](#)

Control reflection

- Which source is authoritative and current?
- Which number, formula or action needs independent review?
- What evidence proves the control ran?
- When must the user refuse, escalate or fall back?

Topic 5: Multi-Agent Orchestration of Financial Work Processes

Supervisor and specialist agents, connected agents, handoff contracts, month-end close orchestration and evaluation

Competency focus: K4 · A3

Core concepts

Why multiple agents

Specialisation buys clarity and costs orchestration.

- Each agent has a narrow scope
- Easier to test and to explain
- Handoffs add new failure modes
- Only split when the split earns its keep

Supervisor and specialists

A supervisor routes; specialists do bounded work.

- Supervisor owns intake and routing
- Specialists own one process each
- The supervisor never becomes an approver
- Escalation always reaches a person

Connected agents

Copilot Studio lets one agent call another published agent.

- Only published agents can be connected
- A description drives the routing choice
- Each agent keeps its own knowledge
- Test the routing, not just the answers

The month-end close crew

Close is the natural first multi-agent finance process.

- A reconciliation specialist
- A variance analyst
- A commentary drafter
- A controller approval gate

Handoff contracts

A handoff needs a structured payload, not a paragraph.

- Task and scope
- Input and source identifiers
- Result and confidence
- Exceptions and required approval

Orchestration patterns

Choose the simplest pattern the process allows.

- Sequential for dependent steps
- Parallel for independent analysis
- Conditional for exception routing
- Escalation for anything unresolved

State and context

Agents lose context unless the design carries it.

- Pass identifiers, not prose
- Persist the run reference
- Keep the source of truth external
- Never rely on conversational memory for evidence

Failure modes

Multi-agent errors compound quietly.

- An error passed on as fact
- Two agents disagreeing silently
- A loop with no termination
- A missing dependency treated as zero

Evaluating the system

Measure the workflow, not only each agent.

- End-to-end task completion
- Groundedness and citation quality
- Handoff success rate
- Control adherence and escalation rate

UAT for a multi-agent process

Test the seams, because that is where it breaks.

- Happy path end to end
- A specialist returning an exception
- An unavailable dependency
- An attempt to bypass the approval gate

Monitoring in production

A live agent system needs owners and telemetry.

- Usage, latency and failure rate
- Escalation and override counts
- Data-policy alerts
- Cost and capacity per process

When not to orchestrate

Some finance work should stay human and single-threaded.

- Material journal approval
- Regulatory interpretation
- Customer eligibility decisions
- Anything irreversible and unbounded

Finance use-case catalogue

Use case	Operating pattern	Value	Risk / control
Month-end close crew	Supervisor routes to reconciliation, variance and commentary	Predictable close	Compounding agent errors / Structured handoffs and a controller gate
Reconciliation specialist	Classify matches, tolerances and exceptions	Consistent treatment	Silent auto-clearing / Policy-defined tolerances and reviewer sign-off
Variance specialist	Identify material variances and draft commentary	Faster analysis	Invented drivers / Causal claims labelled as hypotheses
Collections specialist	Rank overdue accounts and draft outreach	Improved DSO	Unapproved external contact / Draft-only output and human approval
Supervisor routing design	Route by description, not by keyword luck	Predictable behaviour	Wrong specialist chosen silently / Routing tests as part of UAT
Handoff contracts	Pass task, scope, source IDs, result and exceptions	Traceable analysis	Context lost between agents / Structured payloads, never prose
Disagreement surfacing	Show conflicting specialist answers to a human	Better decisions	False consensus / Supervisor never picks a winner silently
Multi-agent UAT	Test the seams, not just the happy path	Fewer production failures	Untested failure modes / Missing data, tool failure and bypass tests

Real-world cases

Balance-sheet reconciliation agent

An agent pattern accelerates matching, exception analysis and close evidence without removing reviewer sign-off.

Source: [Close](#)

The specialist that refused

Asked to rank overdue customers and draft a reminder, the Collections specialist could not retrieve the data and refused to draft on placeholder figures, because a reminder built on guessed numbers reaches a real customer. The supervisor reported the block instead of inventing an answer.

Source: [Course tenant](#)

Control reflection

- Which source is authoritative and current?
- Which number, formula or action needs independent review?
- What evidence proves the control ran?
- When must the user refuse, escalate or fall back?

Topic 6: AI Security, PDPA and Human Oversight in Finance

PDPA obligations, data leakage, data ownership, human decision rights, oversight models and Copilot Studio security settings

Competency focus: K3 · A4 · A5

Core concepts

The finance AI risk picture

Risk spans data, model, access, action and accountability.

- Data leakage and oversharing
- Hallucination and unsupported claims
- Prompt injection and tool abuse
- Unclear accountability for outcomes

PDPA in the AI context

Singapore's PDPA applies to personal data however it is processed.

- Consent, notification and purpose limitation
- Accuracy and protection obligations
- Retention limitation and disposal
- Accountability with a named DPO

PDPA and finance data

Finance holds some of the most sensitive personal data an organisation has.

- Payroll and employee records
- Customer accounts and transactions
- Credit and collections history
- Never paste these into an ungoverned tool

Purpose limitation for AI

Data collected for one purpose is not automatically available for AI.

- Check the original collection purpose
- Assess whether AI use is consistent
- Notify or obtain consent where required
- Document the assessment

Data leakage paths

Leakage is usually an access mistake, not an exotic attack.

- Oversharing through broad site permissions
- Copying sensitive data into prompts
- Agents published without authentication
- Connectors reaching outside the tenant

Data ownership

Ownership answers who decides, who is accountable, and who bears the risk.

- The business owns the data and its use
- IT operates the platform
- Vendor terms govern provider processing
- Ownership must be written down

Vendor terms and data use

Read what the provider may do with your prompts and outputs.

- Whether inputs train the model
- Where data is stored and processed
- Retention and deletion terms
- Incident notification obligations

Prompt injection

Untrusted content can redirect an agent that trusts what it reads.

- Treat retrieved content as data, not instruction
- Segment and filter knowledge sources
- Constrain tools independently of the prompt
- Test adversarial inputs before publishing

The human role in decisions

Some decisions must remain human regardless of model quality.

- Material or irreversible financial decisions
- Decisions affecting a customer's rights
- Regulatory interpretation and disclosure
- Anything requiring delegated authority

Human oversight models

Oversight is a design choice with three distinct shapes.

- Human-in-the-loop approves each action
- Human-on-the-loop monitors and intervenes
- Human-in-command sets and withdraws authority
- Match the model to the impact

Making oversight real

Oversight fails when the reviewer cannot realistically disagree.

- Give the reviewer the evidence
- Allow enough time to review
- Make rejection as easy as approval
- Measure the override rate

Copilot Studio authentication

Authentication decides who the agent will speak to at all.

- Authenticate with Microsoft for internal agents
- No authentication exposes a demo website
- Never use no-auth for finance knowledge
- Least privilege on the connection identity

Data loss prevention policies

DLP separates connectors into business, non-business and blocked.

- Group finance connectors as business
- Block consumer storage and social
- Prevent risky connector combinations
- Policies enforce at run time

Copilot Studio security settings

Several settings decide the real exposure of a published agent.

- Authentication mode
- Channel availability and web publishing
- Knowledge source scope
- Tool and connector permissions

Governance that lasts

Reassess when data, models, regulation or context change.

- Maintain an agent inventory with owners
- Periodic access and content review
- Revalidation triggers on change
- Retire agents nobody uses

Finance use-case catalogue

Use case	Operating pattern	Value	Risk / control
PDPA compliance review	Check purpose, consent and minimisation for AI use	Lawful processing	Purpose creep / Documented assessment and DPO oversight
Data classification for AI	Decide what may and may not enter a prompt	Reduced leakage	Restricted data in prompts / Four-tier classification and prohibitions
Oversharing detection	Find content an agent can reach but should not	Reduced exposure	Broad site permissions / Permission review before publishing an agent
Prompt-injection testing	Run adversarial inputs against a finance agent	Stronger resilience	Testing against live data / Isolated environment and synthetic data only
DLP for connectors	Group connectors as business, non-business or blocked	Prevented exfiltration	Risky connector combinations / Policies enforced at run time
Authentication settings	Require Microsoft sign-in for internal	No anonymous access	Public agent over finance data /

	agents		Authenticate with Microsoft, never no-auth
Approval gate hardening	Make inaction fail safe, not auto-approve	Real oversight	Rubber-stamp approvals / Blank default and reject-on-inaction
Agent inventory and review	Register owner, purpose, scope and review date	Governable estate	Agent sprawl / Annual review and retirement of unused agents
Cost and capacity control	Track credit consumption per process	Predictable run cost	Runaway spend / Usage budgets and environment-level plans
Incident response for AI	Define what happens when an agent gets it wrong	Faster containment	No owner at the moment of failure / Named owner, runbook and rollback

Real-world cases

PDPC guidance on AI systems

Singapore's advisory guidelines address consent, notification, purpose limitation and accountability when personal data is used in AI recommendation and decision systems.

Source: [Regulation](#)

MAS FEAT principles

Fairness, Ethics, Accountability and Transparency give Singapore financial institutions a practical governance lens for AI-assisted decisions.

Source: [Financial services](#)

IMDA Model AI Governance Framework for GenAI

The framework sets out accountability, data quality, testing, incident reporting and content provenance dimensions for generative AI deployment.

Source: [Governance](#)

Control reflection

- Which source is authoritative and current?
- Which number, formula or action needs independent review?
- What evidence proves the control ran?
- When must the user refuse, escalate or fall back?

Your lab accounts and resources

All finance data is synthetic

The Northstar Finance scenario, its policies, reports and lists are created for training. Never substitute live customer, payroll or account data.

Sign-in accounts

Purpose	Sign in with	Password	Portal
Microsoft 365 Premium (Copilot in Word, Excel, PowerPoint)	training1-tertiary (outlook.com)	Trainer will provide	m365.cloud.microsoft
Microsoft 365 Premium (second group)	training2-tertiary (outlook.com)	Trainer will provide	m365.cloud.microsoft
Tenant account (SharePoint, Copilot Studio, agents)	training1 (tenant account)	Trainer will provide	m365.cloud.microsoft
Tenant account (second group)	training2 (tenant account)	Trainer will provide	m365.cloud.microsoft
Trainer and admin (site and environment owner)	admin (tenant account)	Trainer will provide	m365.cloud.microsoft

Which account when

Your trainer will give you the full sign-in details in class. The Premium accounts carry the Copilot 365 licence used for Word, Excel and PowerPoint on Days 1 and 2. The tenant accounts are used for SharePoint, Copilot Studio and the agents from Day 3.

Where the lab data lives

Resource	What is there	Used in
Northstar Finance SharePoint site	Finance Policies and Finance Reports libraries plus seven finance lists	Topics 3 to 6
Finance Policies library	Expense and claims, close calendar, delegation of authority, AI usage policy	Agent grounding
Finance Reports library	P&L, AR ageing, bank reconciliation, forecast and pipeline reports	Agent grounding
Course Power Platform environment	TGS-2026065050 sandbox with Dataverse, agents and	Topics 4 and 5

	agent flows	
labs/_data in the course pack	The same finance data as CSV for the Excel labs	Topic 2

Control totals

Every artefact in this course reconciles to the same figures. If your workbook, list or report disagrees with these, something is wrong.

Measure	Value (SGD)
General ledger revenue	30,402,639.18
General ledger cost	(25,635,379.80)
Net profit H1 FY2026	4,767,259.38
Accounts receivable outstanding at 30 June 2026	8,964,691.11

The four approved policies

Document	ID	Effective	Owner
Expense and Claims Policy	POL-FIN-001	1 January 2026	Financial Controller
Month-End Close Calendar and Controls	POL-FIN-002	1 January 2026	Financial Controller
Delegation of Authority	POL-FIN-003	1 March 2026	Chief Financial Officer
AI Usage and Data Handling Policy	POL-FIN-004	1 May 2026	CFO and Data Protection Officer

Thresholds you will use repeatedly

Expense approval: line manager to SGD 250, department head to SGD 2,000, Financial Controller to SGD 10,000, CFO above. Claims are due within 30 days. Materiality is SGD 50,000 or 10% of budget, whichever is lower. Reconciliation tolerance is SGD 50. Write-off: Financial Controller to SGD 25,000, CFO above.

Detailed lab procedures

Connected Northstar Finance scenario

All twelve labs use one synthetic organisation. Learners progress from classifying AI types, through Copilot in Word, PowerPoint and Excel, to a governed SharePoint site, grounded agents, workflows with human approval, multi-agent orchestration, and a PDPA and security review.

Lab 1: Compare Generative, Agentic and Agent AI for Finance

Goal

Classify finance and fintech use cases by AI type and match each to the control it needs.

Field	Requirement
Topic / criteria	1 / K2 · A1
Duration	60 minutes
Scenario	Northstar Finance is deciding where to start with AI. The CFO wants to know which of the three AI types each candidate use case actually needs, because the controls differ.
Output	A classified use-case map with the control and accountable owner for each.

Before you begin

- Open labs/01-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Review the twenty candidate use cases in the workbook and the definitions sheet.

- Open the starter workbook and read the Definitions sheet before you classify anything.
- The three definitions are not interchangeable labels. Generative produces content, Agentic pursues a goal across steps, and an Agent is the deployed unit with an identity and permissions.
- Note the third column: what the control must address changes with each type.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Classify each as Generative, Agentic or Agent, and record why.

- Work down the Use Case Map sheet one row at a time.
- Ask: does this only produce content (Generative), does it take several steps toward a goal (Agentic), or is it a named deployment with its own permissions (Agent)?
- Record your reason in the same row. A classification without a reason cannot be challenged later.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: For each, state the highest-impact action the AI could take without a person.

- For each row, write the worst thing the AI could do if no person intervened.
- Be specific. 'Make a mistake' is not an answer; 'post a journal that misstates the accounts' is.
- This column is what drives the control, so do not skip it.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Assign the control that must exist before the use case may run.

- Name the control that must exist before the use case may run.
- Prefer structural controls (an approval gate, a permission boundary) over procedural ones (an instruction in a prompt).
- Seven of the twenty rows are decisions that must never be delegated to AI at all. Find them.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Name an accountable human owner for each use case.

- Give every row a named accountable human owner, by role.
- If you cannot name an owner, the use case is not ready to start.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Select three to sequence first and justify the order on control readiness, not enthusiasm.

- On the Sequencing sheet, choose three to start.
- Rank on control readiness, not on how interesting the use case is.
- Be ready to defend why the ones you excluded are not yet ready.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Every use case has a type, a justification, a highest-impact action, a named control and an accountable owner; the three selected are justified on control readiness.

Test it

Prove the lab worked before you move on: Every use case has a type, a justification, a highest-impact action, a named control and an accountable owner; the three selected are justified on control readiness. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 2: Draft a Financial Report with Copilot in Word

Goal

Use Copilot in Word to draft a management report from approved figures and verify every number.

Field	Requirement
Topic / criteria	1 / K2 · A1
Duration	90 minutes
Scenario	The H1 FY2026 results are approved. The CFO needs a management report by tomorrow, drafted from the approved figures and nothing else.
Output	A Word management report with every figure traced to the source workbook.

Before you begin

- Open labs/02-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Open the source workbook and confirm the approved control totals before you draft anything.

- Open the starter workbook and go to the Control Totals sheet.
- Note the net profit for H1 FY2026: SGD 4,767,259.38. Every figure you publish must reconcile to this.
- Read the Approved Figures and By Department sheets. These are the only numbers you may use.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Sign in to Microsoft 365 with the Copilot licence and open a blank Word document.

- Go to m365.cloud.microsoft and sign in with the Microsoft 365 Premium account from your account card.
- Confirm the Copilot icon is present in Word. If it is not, you are signed in with the wrong account.
- Open a blank Word document.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Use Copilot to draft the report, referencing the source workbook and stating the reporting period.

- Open Copilot in Word and choose to draft with it.
- Reference the approved workbook so Copilot draws from it rather than inventing figures.
- State the reporting period explicitly: January to June 2026, as at 30 June 2026.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Direct the draft with a prompt that names the role, the approved source, the required sections and the tone.

- Use the prompt on the Prompt sheet of the workbook rather than a one-line request.
- It names the role, the task, the permitted source, the required sections, the controls and the output shape.
- A one-line prompt produces a plausible report. A prompt contract produces a checkable one.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Check every figure in the draft against the workbook and correct any that do not match.

- Open the Verification Log sheet and list every figure that appears in the draft.
- Check each against the Approved Figures sheet and record match or no match.
- Copilot rounds. A rounded figure in a management report is a wrong figure: replace it with the exact one.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Remove or label any statement of cause that the data does not support.

- Read the draft for statements of cause, not just for numbers.
- The workbook shows what changed. It does not show why. Any 'driven by' or 'due to' claim is a hypothesis.
- Either delete the claim or label it explicitly as a hypothesis to confirm with the department head.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Save the approved report; it becomes the source for the PowerPoint lab.

- Save the corrected report. This exact file is the source for the PowerPoint lab.
- If the report is wrong, the deck generated from it will be wrong in the same way and twice as visible.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Every figure in the report matches the source workbook exactly; causal claims are evidenced or labelled as hypotheses; the report names its period and as-at date.

Test it

Prove the lab worked before you move on: Every figure in the report matches the source workbook exactly; causal claims are evidenced or labelled as hypotheses; the report names its period and as-at date. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out


- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 3: Build a Financial Presentation and an M365 Finance Agent

Goal

Generate a board deck from the approved report, then build a no-code M365 agent grounded on the finance site.

Field	Requirement
Topic / criteria	1 / K2 · A1
Duration	90 minutes
Scenario	The approved report exists. The CFO needs a short board deck from it, and finance staff keep asking the same policy questions that an agent could answer.
Output	A board deck generated from the approved report and a working M365 Copilot agent.



The Northstar Finance site your agent is grounded on.

Before you begin

- Open labs/03-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Open PowerPoint and use Copilot to create a presentation from the approved Word report.

- Open PowerPoint with the same Microsoft 365 Premium account and start a new presentation.
- Use Copilot to create a presentation from the Word report you approved in the previous lab.
- Point it at the file. Do not describe the report from memory.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Correct the story: lead with the decision, not the chronology, and cut slides that repeat the report.

- Compare what Copilot produced against the Deck Story sheet.
- Copilot orders slides the way the document runs. A board deck leads with the decision, not the chronology.
- Move the ask to slide one, and delete slides that only repeat the report.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Cross-check every repeated figure against the report and the source workbook.

- Fill in the Deck Checklist sheet: every figure that appears on a slide, and its value in the report.
- Numbers drift between a document and a generated deck. This is exactly where it happens.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Go to the SharePoint Build page and create an agent, naming it and stating its purpose.

- Go to the SharePoint start page and open the Build tab, then choose Agent.
- Give the agent a clear name and write a purpose that states what it answers and for whom.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Add the Northstar Finance site as the source, so the agent answers only from approved content.

- On the Sources tab, search for the Northstar Finance site by name and add the entire site.
- If the site does not appear, it is not yet in the search index. A newly created site takes 30 to 60 minutes.
- Adding the whole site is right here because the agent needs both the policies and the reports.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Write instructions that require citations, forbid guessing and state that the agent approves nothing.

- On the Behavior tab, write the agent instructions.
- Require it to name the source document and its effective date in every answer.
- Require it to refuse when the approved documents do not cover the question, and to point to Finance Operations.
- State that the agent approves nothing and must name the approver required instead.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Add three starter prompts, create the agent, and test a policy question with a known answer.

- Add three starter prompts so a new user can see what the agent is for.
- Create the agent, then open it and ask a question whose answer you already know.

- Check the citation, not just the answer. An answer without a source is not usable finance evidence.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Test a question the documents cannot answer and confirm the agent refuses rather than inventing.

- Now ask something the documents cannot answer, such as a cryptocurrency treasury policy.
- A correct agent says it cannot find that in the approved documents and directs you elsewhere.
- An agent that invents a plausible answer has failed the test, however good the answer sounds.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Repeated figures reconcile across workbook, report and deck; the agent cites a source document and its effective date, and refuses an out-of-scope question instead of guessing.

Test it

Prove the lab worked before you move on: Repeated figures reconcile across workbook, report and deck; the agent cites a source document and its effective date, and refuses an out-of-scope question instead of guessing. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 4: Prepare Finance Data for Copilot in Excel

Goal

Convert a messy finance export into an analysis-ready Excel Table with reconciled control totals.

Field	Requirement
Topic / criteria	2 / K1 · A2
Duration	90 minutes
Scenario	The FP&A team receives a mixed-quality export from the billing system. Copilot will answer confidently on dirty data, so the data must be fixed first.
Output	A clean transactions table, a quality log and reconciled control totals.

Before you begin

- Open labs/04-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Open the starter workbook and read the data dictionary and the expected control total.

- Open the starter workbook and read the Data Dictionary sheet first.
- It states the expected type and the rule for every column. That is your definition of clean.
- Note the control totals on the Quality and Controls sheet before you change anything.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Ask Copilot to describe the data quality issues, then verify each claim yourself.

- Open Copilot in Excel and ask it to describe the data quality issues in the Transactions sheet.
- Record what it says. Then verify each claim yourself against the actual rows.
- Copilot will answer confidently on dirty data. The point of this step is to see that happen.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Convert the range to an Excel Table with Ctrl+T and standardise the headers.

- Select the data range and press Ctrl+T, confirming that the table has headers.
- Give the table a name in Table Design. This is the single highest-value action before using Copilot.
- A named table gives Copilot an explicit grain and column set instead of a guess.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Correct dates, currencies and account codes without altering any transaction ID.

- Fix the seeded defects: a blank department, an invalid date, a duplicate journal ID, a blank account code, an amount stored as text, a cost with a positive sign, an impossible period and a blank amount.
- Never edit a JournalID to resolve a duplicate. Investigate which row is wrong.
- Convert text-formatted numbers to real numbers, or they will not sum.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Add validation columns for missing keys, duplicates and sign errors.

- Build the three validation columns: required fields, duplicate check and sign check.
- The sign rule is that 4xxx accounts are positive revenue and everything else is negative cost.
- Conditional formatting will highlight the failures as you go.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Reconcile the cleaned revenue and expense totals to the supplied control figures.

- Reconcile the cleaned revenue and cost totals to the Quality and Controls sheet.
- Net profit must equal SGD 4,767,259.38. If it does not, a row is still wrong or has been lost.
- Row count must be unchanged. Cleaning never deletes rows.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Repeat your first Copilot question on the cleaned table and compare the two answers.

- Ask Copilot your original question again, now on the clean table.
- Record both answers side by side on the Copilot Comparison sheet.
- The gap between them is the lesson: the tool did not change, the data did.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Duplicate count is zero, required-field errors are zero, cleaned totals reconcile to the control figures, and the before-and-after Copilot answers are documented.

Test it

Prove the lab worked before you move on: Duplicate count is zero, required-field errors are zero, cleaned totals reconcile to the control figures, and the before-and-after Copilot answers are documented. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 5: Analyse Variances and Build a Dashboard with Copilot

Goal

Use Copilot in Excel to analyse budget variances and build an auditable management dashboard.

Field	Requirement
Topic / criteria	2 / K1 · A2
Duration	150 minutes
Scenario	Northstar's CFO needs a monthly performance dashboard and commentary that survives scrutiny in the board meeting.
Output	An editable dashboard with actual, budget, variance and forecast views plus verified commentary.

Before you begin

- Open labs/05-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Review the Actuals, Budget and Assumptions tables and confirm they are proper Excel Tables.

- Open the workbook and confirm the Actuals, Budget and Assumptions ranges are proper Excel Tables.
- If any is a plain range, convert it with Ctrl+T before continuing.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Ask Copilot to generate the SUMIFS, XLOOKUP and IFERROR formulas in the Analysis sheet.

- Ask Copilot to build the department summary using SUMIFS against the PL Data table.
- Ask for XLOOKUP where you need a mapped dimension, and IFERROR to handle blanks explicitly.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Audit every generated formula: check ranges, absolute references and blank handling.

- Audit every generated formula before you trust it.
- Check the ranges cover the whole table, check absolute references are anchored, and test what happens on a blank.
- A generated formula is a proposal. It deserves the same audit as one a colleague wrote.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Apply the materiality rule from the close policy and identify the material variances.

- Apply the materiality rule from POL-FIN-002: greater than SGD 50,000 or 10% of budget, whichever is lower.
- Flag the material lines. Aggregate the rest rather than listing them.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Ask Copilot to create charts, then choose the chart type that matches the claim you are making.

- Ask Copilot to create charts, then decide whether the chart type matches the claim.
- Columns compare periods, lines show trend, and a waterfall shows a bridge. A pie chart shows almost nothing useful in finance.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Build KPI tiles linked to the Analysis table, never to typed-in values.

- Build the KPI tiles on the Dashboard sheet as formulas referencing the Analysis table.
- Never type a value into a tile. A typed tile is correct exactly once, and then silently wrong forever.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Ask Copilot to draft commentary, then delete every causal claim the data does not evidence.

- Ask Copilot to draft the commentary, then read it as a sceptic.
- Delete every causal claim the data does not evidence, or relabel it as a hypothesis with an owner to confirm it.
- Record the result on the Commentary sheet, separating fact from hypothesis in the dedicated column.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Run the formula error scan and reconcile the dashboard to the control totals.

- Run a formula error scan across the workbook and resolve every error.
- Reconcile the dashboard KPI tiles to the control totals. The Reconciles column must read OK on every row.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

All formula checks pass; total variance equals actual less budget; charts stay linked to source data; commentary separates evidenced fact from labelled hypothesis.

Test it

Prove the lab worked before you move on: All formula checks pass; total variance equals actual less budget; charts stay linked to source data; commentary separates evidenced fact from labelled

hypothesis. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 6: Reconcile Accounts and Forecast with Copilot

Goal

Apply deterministic matching rules and build a driver-based forecast with scenarios.

Field	Requirement
Topic / criteria	2 / K1 · A2
Duration	120 minutes
Scenario	The controller must reconcile the bank statement to the receivables ledger and produce a defensible H2 forecast.
Output	A reconciliation report with an owned exception queue and a base, downside and upside forecast.

Before you begin

- Open labs/06-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Review the bank and receivables tables and the match rules from the close policy.

- Read the Match Rules sheet. The four classifications come from POL-FIN-002 and are not negotiable.
- Note the escalation rule: above SGD 10,000 or older than 60 days goes to the Financial Controller.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Apply exact, tolerance and reference-only match logic in the matching sheet.

- Work through the Bank Statement sheet and classify each line.
- Exact means amount and reference both agree. Tolerance means the amount differs by SGD 50 or less.
- Reference-only means the amount agrees but the narrative lost the invoice ID; it always needs reviewer confirmation.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Classify every remaining line as an exception, with an age and an owner.

- Everything that is not a match is an exception. Give each one an age and a named owner.
- Bank-only items such as charges, FX adjustments and returned payments are exceptions, not noise.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Confirm matched plus exception values reconcile to each source total.

- On the Reconciliation sheet, confirm that lines classified equals bank lines.
- An unclassified line is not a clean reconciliation; it is an unfinished one.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Build the forecast from assumption cells for volume, price and cost, never hard-coded outputs.

- Move to the Forecast Assumptions sheet and set the driver cells.
- Every forecast figure must reference an assumption cell. A hard-coded output cannot be re-run.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Create base, downside and upside scenarios by changing assumptions coherently.

- Build the base, downside and upside scenarios on the Forecast sheet.
- Change assumptions coherently: a downside that cuts revenue but leaves costs untouched is not a scenario, it is an error.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Ask Copilot to summarise the reconciliation and the forecast, then verify every figure it states.

- Ask Copilot to summarise the reconciliation and the forecast.
- Verify every figure it states against your own sheets before you accept the summary.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Matched plus exception values reconcile to both sources; every unresolved item has a reason, age and owner; scenarios are driven by assumption cells; escalation items above the policy threshold are flagged.

Test it

Prove the lab worked before you move on: Matched plus exception values reconcile to both sources; every unresolved item has a reason, age and owner; scenarios are driven by assumption cells; escalation items above the policy threshold are flagged. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 7: Build the Finance SharePoint Site and Load Finance Data

Goal

Create a governed finance site with typed lists and an approved document library.

Field	Requirement
Topic / criteria	3 / K1 · A3
Duration	120 minutes
Scenario	Finance content is scattered across personal drives. The team needs one governed location that an agent can safely be grounded on.
Output	A finance site with policy and report libraries, typed lists and documented permissions.

The screenshot displays the SharePoint interface for a site named 'Northstar Finance'. The left-hand navigation pane includes options like 'Discover', 'Publish', 'Build', and 'OneDrive'. The 'Site contents' section is active, showing a list of documents and lists. The main area shows a table with the following data:

Name	Type	Items	Modified
Documents	Document library	0	9/5/2026 10:59 PM
Finance Policies	Document library	8	9/6/2026 12:54 AM
Finance Reports	Document library	10	9/6/2026 12:54 AM
Form Templates	Document library	0	9/5/2026 10:49 PM
Site Assets	Document library	4	9/5/2026 10:48 PM
Style Library	Document library	0	8/23/2026 8:56 AM
Bank Statement	List	17	9/5/2026 10:56 PM
Chart of Accounts	List	10	9/5/2026 10:54 PM
Collections AR	List	33	9/5/2026 10:56 PM

Site Contents: two document libraries and seven typed finance lists.

Before you begin

- Open labs/07-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Create a team site named for the finance team and record its URL.

- Sign in to SharePoint with your tenant account.
- From the Build page choose Site, then the Standard team template.
- Name the site for the finance team, set it private, and record the resulting URL on the Site Build sheet.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Create the Finance Policies and Finance Reports document libraries with clear descriptions.

- Create two document libraries: Finance Policies and Finance Reports.
- Give each a description that says what belongs in it. The description is what an agent sees when choosing a source.
- Keep them separate: policy changes rarely and needs version control; reports are point-in-time and are superseded each period.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Upload the supplied policy documents and period reports.

- Upload the four policy documents to Finance Policies and the five period reports to Finance Reports.
- Upload the Word versions. Copilot Studio will not index Markdown, so a .md-only library grounds an agent on nothing.
- Check each document carries its ID, effective date and owner. A policy without an effective date cannot be cited safely.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Create the finance lists and set every column to its correct type: currency, date, number or choice.

- Create the seven lists from the Lists sheet.
- Set every column to its correct type as you create it. Currency for money, Date for dates, Choice for controlled values.
- A number stored as text cannot be summed, filtered or charted, and the mistake is expensive to undo later.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Import the supplied CSV data into the lists and verify the row counts.

- Import the CSV data from the labs/_data folder into the matching lists.
- Check the row counts against the Lists sheet: 240, 240, 240, 140, 33, 17 and 10.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Reconcile the loaded data to the control totals in the data pack.

- Reconcile the loaded data on the Reconciliation sheet.
- Net profit must come to SGD 4,767,259.38 and AR outstanding to SGD 8,964,691.11.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Review the site permissions and record who can see what, and why.

- Open the site permissions and complete the Permissions sheet.
- An agent grounded here can surface anything the asking user may already open. Broad access becomes broad disclosure.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Note which content is sensitive enough to need a sensitivity label.

- Identify which content would need a sensitivity label in a real deployment.
- Nothing restricted is stored on this training site, and that is itself a design decision worth stating.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Both libraries and all seven lists exist with correct column types; row counts match the data pack; loaded data reconciles to the control totals; the permission review names each group and its justification.

Test it

Prove the lab worked before you move on: Both libraries and all seven lists exist with correct column types; row counts match the data pack; loaded data reconciles to the control totals; the permission review names each group and its justification. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 8: Ground a Finance Agent and Test Its Refusals

Goal

Build a Copilot Studio agent grounded on approved finance content and prove it refuses correctly.

Field	Requirement
Topic / criteria	3 / K1 · A3
Duration	120 minutes
Scenario	Finance staff need reliable answers on expenses, close deadlines and delegated authority, without the agent ever inventing a threshold.
Output	A grounded finance agent and a completed UAT log including refusal evidence.

The screenshot displays the Copilot Studio 'Agents' page. On the left, a sidebar contains 'Home', 'Agents' (selected), and 'Workflows'. The main content area is titled 'Agents' with a sub-header '5 items'. Below this is a search bar 'Search by name' and a 'New agent' button. A table lists the agents:

Name	Status	Owner	Powered by
Month End Close Supervisor	Published	TIA Admin	GitHub Copilot
Collections Assistant	Published	TIA Admin	GitHub Copilot
Close Reconciliation Agent	Published	TIA Admin	GitHub Copilot
Finance Policy Assistant	Published	TIA Admin	GitHub Copilot
FPA Variance Analyst	Published	TIA Admin	GitHub Copilot

At the bottom left of the sidebar, the environment name 'TGS-2026065050-Generative A...' and the user 'TIA Admin' are visible.

The agents in the course Copilot Studio environment.

Before you begin

- Open labs/08-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Create an agent in the course environment and give it a clear finance role name.

- Open Copilot Studio and confirm you are in the course environment before you create anything.
- The environment name appears at the bottom left. Building in the wrong environment is a common and annoying mistake.
- Create a new agent and give it a clear finance role name.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Write instructions that name the scope, require citations and effective dates, and forbid guessing.

- Write the instructions as prose, not as bullet fragments.
- Name the scope: the four approved finance policies and nothing else.
- Require a citation and an effective date in every answer.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: State explicitly that the agent approves nothing and must name the required approver instead.

- Add the refusal and authority rules explicitly.
- State that the agent never guesses a threshold, approver or deadline.
- State that the agent approves nothing and names the required approver instead.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Add the Finance Policies library as a SharePoint knowledge source.

- Add the SharePoint knowledge source pointing at the Finance Policies library.
- Paste the URL with spaces percent-encoded as %20. With raw spaces the Add button stays disabled and it is not obvious why.
- Point at the Word versions of the policies. Copilot Studio does not index Markdown files: the source shows Status Ready, and every search returns nothing.
- Microsoft 365 Copilot agents do index Markdown, which is why the same documents can work in one surface and not the other.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Remove the default web search source so the agent cannot answer from the open internet.

- Remove the default Search all websites knowledge chip.
- Left in place, the agent can answer from the open internet and your grounding is decorative rather than real.
- Save the agent first: removing the chip before the first save is silently reverted.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Save, publish, and confirm the agent shows as published before testing anything.

- Publish the agent and wait for the confirmation. Publishing takes 30 to 60 seconds.
- Check the agent list shows Published, not Draft, before you test anything.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Run the positive tests: questions whose answers are in the policies, checking the citation each time.

- Run the four positive cases from the UAT Cases sheet.
- For each, record what the agent cited, not just whether the answer looked right.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Run the negative tests: an out-of-scope question, a superseded policy and a request to approve something.

- Run the negative, authority, injection and data cases.
- These matter more than the positive ones. An agent that answers well but cannot refuse is not safe to deploy.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Step 9: Record every result in the UAT log with the evidence, and retest after any instruction change.

- Record every defect on the Defects sheet with the instruction change you made and the retest result.
- If the agent retrieves nothing at all, read the If It Returns Nothing sheet before assuming it is broken.

Evidence checkpoint

Record completion of step 9, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

The agent cites the correct document and effective date on positive tests; refuses on out-of-scope, unsupported and approval requests; the UAT log records evidence for every case; web search is removed.

Test it

Prove the lab worked before you move on: The agent cites the correct document and effective date on positive tests; refuses on out-of-scope, unsupported and approval requests; the UAT log records evidence for every case; web search is removed. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 9: Create the Environment and a Finance Agent Flow

Goal

Provision a governed environment and build an agent flow that reads approved finance data.

Field	Requirement
Topic / criteria	4 / A3
Duration	120 minutes
Scenario	Course work must be isolated from production. The finance team then wants a flow that assembles the reconciliation position on demand.
Output	A course environment and a published agent flow with documented tools and scope.

Before you begin

- Open labs/09-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Create a sandbox environment named for the course, in the Asia region with SGD currency.

- Open the Power Platform admin centre and create a new environment.
- Name it for the course code so it can be found and governed later.
- Choose Sandbox as the type and Asia as the region.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Enable Dataverse, because agents and flows require it, and record the environment ID.

- Set the currency to SGD, then open Change default settings and enable the Dataverse data store.
- Agents and agent flows require Dataverse. It cannot be added quietly afterwards.
- Set the security group to None for a training cohort. That would be the wrong choice for production finance data.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Confirm the environment is covered by a billing plan, or agents will fail with a credits error.

- Confirm the environment is a target of a billing plan.
- Without one, every agent reply fails with EnforcementUsageCredits and the cause is not visible from inside Copilot Studio.
- Record the environment ID: every Copilot Studio and admin URL contains it.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Create an agent flow and give it a name ending in the do-not-delete marker.

- In Copilot Studio, create a workflow and name it ending in (DO NOT DELETE).
- Choose the trigger When an agent calls the workflow. Only this trigger makes the flow selectable as an agent tool.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Add a SharePoint Get items action pointing at the finance list, and set a filter query.

- Add a SharePoint Get items action.
- Pick the site from the dropdown, never a typed URL, and select the finance list.
- Build the filter query as literal text wrapping a picker token, and normalise the key inside the filter rather than after it.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Insert every dynamic value with the picker rather than typing an expression.

- Insert every dynamic value with the lightning-bolt picker.
- Typed or pasted expressions are escaped by the editor, and a reference to a missing node resolves to empty rather than erroring.
- The node stays green and the flow returns nothing. This is the single most common failure in this designer.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Add the flow as a tool on your agent, publishing the flow first so it becomes selectable.

- Publish the flow, then add it as a tool on your agent.
- Only published flows appear in the tool picker. An unpublished flow is greyed out.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Test the flow, then read the run history node by node to confirm what actually executed.

- Test the flow, then open the run history and read the node outputs.
- A green run does not mean the work happened: a node left in Needs setup is skipped silently.
- Read the upstream node's outputs before proposing a second fix for any failure.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

The environment exists with Dataverse and a billing plan; the flow is published and runs green; the run history shows the SharePoint action returning rows; the flow appears as a tool on the agent.

Test it

Prove the lab worked before you move on: The environment exists with Dataverse and a billing plan; the flow is published and runs green; the run history shows the SharePoint action returning rows; the flow appears as a tool on the agent. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 10: Add a Human Approval Gate to a Finance Workflow

Goal

Build a blocking human-in-the-loop approval gate and prove that inaction fails safe.

Field	Requirement
Topic / criteria	4 / A3
Duration	90 minutes
Scenario	Collections wants to automate reminder drafting. Nothing may reach a customer without a person approving it first.
Output	A workflow with a working approval gate, an audit trail and evidence of both decision paths.

Before you begin

- Open labs/10-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Extend your flow so it drafts a payment reminder from the approved receivables data.

- Extend your flow so it drafts a payment reminder from the receivables data.
- Use the escalated accounts on the Approval Queue sheet as your test cases.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Write the draft to an audit record before the approval gate, so the evidence exists either way.

- Write the draft to an audit record before the approval gate, not after.
- Logging before the gate is what lets someone later ask whether anything reached a customer that nobody reviewed.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Add a Human review node and define its inputs, including an outcome and an approver name.

- Add a Human review node. Define two inputs: an Outcome as Yes/No and an approver name as Text.
- The node will not save with zero inputs, and there is no built-in outcome property to reference.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Leave the outcome default blank so that inattention cannot become approval.

- Leave the Outcome default blank.

- A default of Approve means the field arrives already answered: confirming takes no thought and rejecting takes noticing.
- That converts a gate into a rubber stamp through a setting invisible on the canvas.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Set the channel to Teams and address the request to a resolved directory account.

- Set the Channel to Teams.
- On a live tenant the Outlook channel created the request but never delivered it, to any address.
- For Assigned to, type the full address, wait for the directory lookup and click the resolved suggestion.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Add an If and Else branch on the outcome and build both paths, including the rejection path.

- Add an If and Else on the Outcome, comparing against the literal Yes.
- The Yes/No input is shown as boolean in the picker but publishes the string Yes. Comparing against true never matches.
- Build the rejection branch too. An empty Else means rejecting silently does nothing.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Publish, then run the approval path and record the evidence.

- Publish, then run the approval path and record the evidence on the Both Paths sheet.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Run the rejection path and confirm nothing is sent and the rejection is recorded.

- Run the rejection path. Confirm nothing is sent and the rejection is recorded.
- A gate that has only ever been approved is untested.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Step 9: Compare a blank default with a pre-filled Approve default and explain the control difference.

- Complete the Default Comparison sheet.
- Explain in your own words why a blank default is a control and a pre-filled Approve is not.

Evidence checkpoint

Record completion of step 9, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

The gate blocks until a person responds; both approval and rejection paths are evidenced; the draft is logged before the gate; the default is fail-safe and the learner can explain why.

Test it

Prove the lab worked before you move on: The gate blocks until a person responds; both approval and rejection paths are evidenced; the draft is logged before the gate; the default is fail-safe and the learner can explain why. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 11: Orchestrate a Multi-Agent Month-End Close

Goal

Build a supervisor with connected specialist agents and prove the routing and the escalation.

Field	Requirement
Topic / criteria	5 / K4 · A3
Duration	75 minutes
Scenario	Reconciliation, variance analysis and collections each need different expertise. A supervisor should route the work without ever becoming an approver.
Output	A working supervisor, a handoff contract, routing tests and an escalation path.

Copilot Studio | Month End Close Supervisor | Build | Preview | Evaluate | Monitor | Publish

Month End Close Supervisor

Instructions

You are the Month End Close Supervisor for Northstar Finance. You coordinate the specialist finance agents through the month-end close and hand a consolidated position to the Financial Controller.

Your role is routing, not doing. You decide which specialist should handle a request, pass it on with the information it needs, and assemble the answers. You do not perform reconciliation, variance analysis or collections work yourself.

Who to route to. Send bank and receivable matching, reconciling items and exception ageing to the Close Reconciliation Agent. Send actual versus budget movements, materiality and management commentary to the FPA Variance Analyst. Send overdue accounts, customer ageing and payment reminders to the Collections Assistant. Send questions about policy, deadlines, thresholds and approval authority to the Finance Policy Assistant.

How to hand off. State the task, the period, the department or account in scope, and what a good answer looks like. When a specialist returns an answer, keep its figures and its caveats intact. Never restate a specialist's number more precisely than the specialist gave it.

Close sequence. Follow the close calendar order: reconciliations complete before variance analysis. and variance commentary completes before the

Model | Claude Opus 5

Channels | Define where users can interact with the agent.

Skills | Define behaviors through structured instructions

Tools | Connect the agent to external systems and actions

Knowledge | Provide trusted context to guide decisions

Connected agents | Collaborate across agents to complete work

- Finance Policy Assistant
- Collections Assistant
- FPA Variance Analyst
- Close Reconciliation Agent

Memory | Preview | Remember interactions, workflows and context for improved results.

The supervisor holding four specialists as connected agents.

Before you begin

- Open labs/11-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Build the specialist agents for reconciliation, variance and collections, each with a narrow scope.

- Build the three specialist agents, each with a narrow scope and an explicit prohibition.
- Use the Architecture sheet: reconciliation, variance and collections.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Publish every specialist, because only published agents can be connected.

- Publish every specialist before going further.
- Connected agents accept published agents only. Unpublished ones simply do not appear in the picker.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Create the supervisor agent and write instructions that make routing its job, not doing the work.

- Create the supervisor agent.
- Write instructions that make routing its job. State plainly that it does not perform the specialist work itself.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Add each specialist as a connected agent with a routing description that says when to use it.

- Add each specialist as a connected agent.
- The description is required and it is what drives routing, so say when to use this specialist, not what it is.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: State the close sequence in the supervisor instructions so a blocked step is reported, not skipped.

- State the close sequence in the supervisor instructions.
- Reconciliation completes before variance analysis; commentary completes before controller review.
- If an earlier step is unresolved the supervisor must report the later step as blocked, not proceed without it.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Complete the handoff contract: task, scope, source IDs, result, confidence and exceptions.

- Complete the Handoff Contract sheet.
- A handoff needs a structured payload. Prose loses precision at every hop between agents.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Test routing with one question per specialist and confirm each reaches the right one.

- Run routing tests R01 to R04, one per specialist.
- Record which specialist actually answered. Routing is driven by your descriptions, so a miss is a description problem.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Test a question that spans two specialists and check the supervisor keeps both answers intact.

- Run R05 and R06, which span two specialists.
- Check the supervisor keeps both sets of figures and both sets of caveats intact.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Step 9: Test a failure case where data is missing and confirm the supervisor reports the block honestly.

- Run R08, where the source data is unavailable.
- The correct behaviour is to report the block and name the missing input, never to produce a plausible number.

Evidence checkpoint

Record completion of step 9, the evidence location, reviewer status and any unresolved exception.

Step 10: Confirm no agent in the system can approve, post or send anything.

- Run R07 and confirm no agent in the system will approve, post or send anything.
- Complete the Failure Modes sheet with what you observed.

Evidence checkpoint

Record completion of step 10, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Each routing test reaches the correct specialist; the supervisor preserves specialist figures and caveats; a missing-data case is reported as blocked rather than answered; no agent can approve, post or send.

Test it

Prove the lab worked before you move on: Each routing test reaches the correct specialist; the supervisor preserves specialist figures and caveats; a missing-data case is reported as blocked rather than answered; no agent can approve, post or send. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Lab 12: Secure a Finance Agent for PDPA and Human Oversight

Goal

Apply PDPA, data-classification, security and oversight controls, then decide whether the agent may be published.

Field	Requirement
Topic / criteria	6 / K3 · A4 · A5
Duration	60 minutes
Scenario	The finance agent is proposed for wider deployment. Before it is released, it must pass a PDPA review, a security review and an oversight review.
Output	A completed security and PDPA review with adversarial test evidence and a publish decision.

Agents
5 items

Search by name

New agent

Name	Status	Owner	Powered by
Month End Close Supervisor	Published	TIA Admin	GitHub Copilot
Collections Assistant	Published	TIA Admin	GitHub Copilot
Close Reconciliation Agent	Published	TIA Admin	GitHub Copilot
Finance Policy Assistant	Published	TIA Admin	GitHub Copilot
FPA Variance Analyst	Published	TIA Admin	GitHub Copilot

Every agent needs a named owner, a scope and a review date.

Before you begin

- Open labs/12-*/README.pdf and the starter workbook.
- Work only with the supplied synthetic data; never paste live finance or customer data.
- Save a working copy with your name or assigned learner ID.
- Keep formulas, prompts, screenshots and review notes as evidence.

Step 1: Inventory the agent: purpose, owner, users, knowledge sources, tools and connectors.

- Complete the Agent Inventory sheet for the agent you are reviewing.
- You cannot govern an agent that has no recorded owner, purpose or scope.

Evidence checkpoint

Record completion of step 1, the evidence location, reviewer status and any unresolved exception.

Step 2: Complete the PDPA review: purpose, consent basis, minimisation, retention and the DPO contact.

- Work through the PDPA Review sheet obligation by obligation.
- Purpose limitation is the one most often failed: data collected for billing is not automatically available for AI.

Evidence checkpoint

Record completion of step 2, the evidence location, reviewer status and any unresolved exception.

Step 3: Classify every data source the agent can reach and confirm no restricted data is in scope.

- Classify every data source the agent can reach on the Data Classification sheet.
- Restricted data, meaning payroll, NRIC, bank details and credit files, must never be in scope for any AI prompt.

Evidence checkpoint

Record completion of step 3, the evidence location, reviewer status and any unresolved exception.

Step 4: Check the authentication setting and justify it; internal finance content must never be no-auth.

- Open the agent's Settings and check the authentication mode.
- Authenticate with Microsoft is correct for internal finance content.
- No authentication publishes a demo website reachable by anyone with the link. Justify your setting in writing.

Evidence checkpoint

Record completion of step 4, the evidence location, reviewer status and any unresolved exception.

Step 5: Review the knowledge scope and test whether the agent can reach content the asker should not see.

- Review the knowledge scope. Test whether the agent surfaces anything the asking user should not see.
- The narrowest scope that answers the question is the right one.

Evidence checkpoint

Record completion of step 5, the evidence location, reviewer status and any unresolved exception.

Step 6: Run the supplied adversarial prompts, including prompt injection and approval-bypass attempts.

- Run the adversarial prompts from adversarial-prompts.json against the training agent only.
- Use synthetic data. Never run these against production or against real customer records.
- Record the actual response for each, not just pass or fail.

Evidence checkpoint

Record completion of step 6, the evidence location, reviewer status and any unresolved exception.

Step 7: Record which controls are structural and which are only procedural, and rank them.

- Complete the Control Ranking sheet.

- Separate structural controls the model cannot reach from procedural ones written into a prompt, and from conventions such as a self-declared name.

Evidence checkpoint

Record completion of step 7, the evidence location, reviewer status and any unresolved exception.

Step 8: Review the approval gate default and confirm inaction fails safe.

- Check the approval gate default on any connected workflow.
- Confirm that inaction results in rejection rather than approval.

Evidence checkpoint

Record completion of step 8, the evidence location, reviewer status and any unresolved exception.

Step 9: Set out the oversight model and confirm the reviewer has evidence, time and a real ability to reject.

- Complete the Human Oversight sheet.
- Ask honestly whether the reviewer has the evidence, the time and a realistic ability to say no. If not, the oversight is decorative.

Evidence checkpoint

Record completion of step 9, the evidence location, reviewer status and any unresolved exception.

Step 10: Issue a publish, remediate or reject recommendation with residual risk and a named owner.

- Issue your recommendation on the Publish Decision sheet.
- Name the highest residual risk, who accepts it, the review date and the rollback plan.
- A recommendation without a named risk acceptor is not a decision.

Evidence checkpoint

Record completion of step 10, the evidence location, reviewer status and any unresolved exception.

Acceptance criteria

Every high risk has a preventive and a detective control with test evidence; the PDPA review is complete with a named DPO; no restricted data is in scope; adversarial tests are evidenced; the recommendation names a residual risk and an accountable owner.

Test it

Prove the lab worked before you move on: Every high risk has a preventive and a detective control with test evidence; the PDPA review is complete with a named DPO; no restricted data is in scope; adversarial tests are evidenced; the recommendation names a residual risk and an accountable owner. If any part of that is not true, the lab is not finished, whatever the screen shows.

Close-out

- Compare the starter and solution structures only after completing your own work.
- Submit the learner-created output specified by the trainer.
- Remove personal data, credentials and unapproved external links before submission.

Finance AI security checklist

Control area	Questions to answer
Identity and access	Is the user authenticated? Are source permissions preserved? Is least privilege applied?
Data protection	Is financial data classified? Are DLP policies and approved connectors enforced?
Model and prompt risk	Are grounding, refusal, prompt-injection and hallucination tests documented?
Action safety	Are write tools, monetary thresholds and approvals bounded?
Monitoring	Are logs, owners, incidents, cost and drift reviewed?
Accountability	Can the reviewer reproduce the evidence and explain the final decision?

References

[Course page](#)

[Copilot in Excel - Get started](#)

[Copilot in Word](#)

[Copilot in PowerPoint](#)

[Create an agent in SharePoint](#)

[Microsoft 365 Copilot agents overview](#)

[Copilot Studio documentation](#)

[Copilot Studio agent flows](#)

[Copilot Studio connected agents](#)

[Copilot Studio security and governance](#)

[Copilot Studio data loss prevention](#)

[Copilot Studio authentication](#)

[Copilot Studio agent evaluation](#)

[Power Platform environments](#)

[Microsoft Copilot Control System](#)

[SharePoint sensitivity labels](#)

[PDPC - Personal Data Protection Act overview](#)

[PDPC - Advisory Guidelines on use of Personal Data in AI systems](#)

[IMDA Model AI Governance Framework for Generative AI](#)

[MAS FEAT principles](#)

[Microsoft finance scenario library](#)

[Finance Agent overview](#)